

LIVER DISEASE PREDICTION DETECTION

importing the libraries

In [1]:

```
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler

from sklearn.model_selection import train_test_split

from sklearn.model_selection import KFold

from sklearn.model_selection import cross_val_score

from sklearn.model_selection import GridSearchCV

from sklearn.neighbors import KNeighborsClassifier

from sklearn.linear_model import LogisticRegression

from sklearn.tree import DecisionTreeClassifier

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

loading the dataset

In [2]:

```
dataset = pd.read_csv(
    "/home/dhanoia/RUBUNTU/DATASETS/CSV_DATASETS/indian_liver_patient.csv")

dataset.head()
```

Out[2]:

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphatase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Proteins
0	65	Female	0.7	0.1	187	16	18	6
1	62	Male	10.9	5.5	699	64	100	7
2	62	Male	7.3	4.1	490	60	68	7
3	58	Male	1.0	0.4	182	14	20	6
4	72	Male	3.9	2.0	195	27	59	7

checking datasets for null and type mismatch

In [3]:

```
dataset.info()

# Output:
RangeIndex: 583 entries, 0 to 582
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  --
0   Age                    583 non-null    int64
1   Gender                 583 non-null    object
2   Total_Bilirubin        583 non-null    float64
3   Direct_Bilirubin       583 non-null    float64
4   Alkaline_Phosphatase   583 non-null    int64
5   Alamine_Aminotransferase 583 non-null    int64
6   Aspartate_Aminotransferase 583 non-null    int64
7   Total_Proteins         583 non-null    float64
8   Albumin                583 non-null    float64
9   Albumin_and_Globulin_Ratio 579 non-null    float64
10  Dataset                583 non-null    int64
Dtypes: float64(5), int64(5), object(1)
memory usage: 35.97 KB
```

preprocessing the dataset

In [4]:

```
# filling the null values of 'Albumin_and_Globulin_Ratio' column with its mean

dataset['Albumin_and_Globulin_Ratio'].fillna(
    dataset['Albumin_and_Globulin_Ratio'].mean(), inplace=True)

dataset.info()
```

```
RangeIndex: 583 entries, 0 to 582
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  --
0   Age                    583 non-null    int64
1   Gender                 583 non-null    object
2   Total_Bilirubin        583 non-null    float64
3   Direct_Bilirubin       583 non-null    float64
4   Alkaline_Phosphatase   583 non-null    int64
5   Alamine_Aminotransferase 583 non-null    int64
6   Aspartate_Aminotransferase 583 non-null    int64
7   Total_Proteins         583 non-null    float64
8   Albumin                583 non-null    float64
9   Albumin_and_Globulin_Ratio 583 non-null    float64
Dtypes: float64(5), int64(5), object(1)
memory usage: 35.97 KB
```

converting each non-numeric data to numeric for visualization

In [5]:

```
dataset = pd.get_dummies(dataset)
```

visualizing each attribute

In [6]:

```
fig, axes = plt.subplots(nrows=rows, ncols=cols, figsize=(15, 15))
names = dataset.columns.values
counter = 0


for i in range(rows):
    for j in range(cols):
        if counter >= len(names):
            break

        name = names[counter]
        sns.distplot(a=dataset[name], ax=axes[i, j])

        counter += 1
```

In [7]:

```
draw_hist(dataset, dataset, rows=5, cols=5)
```



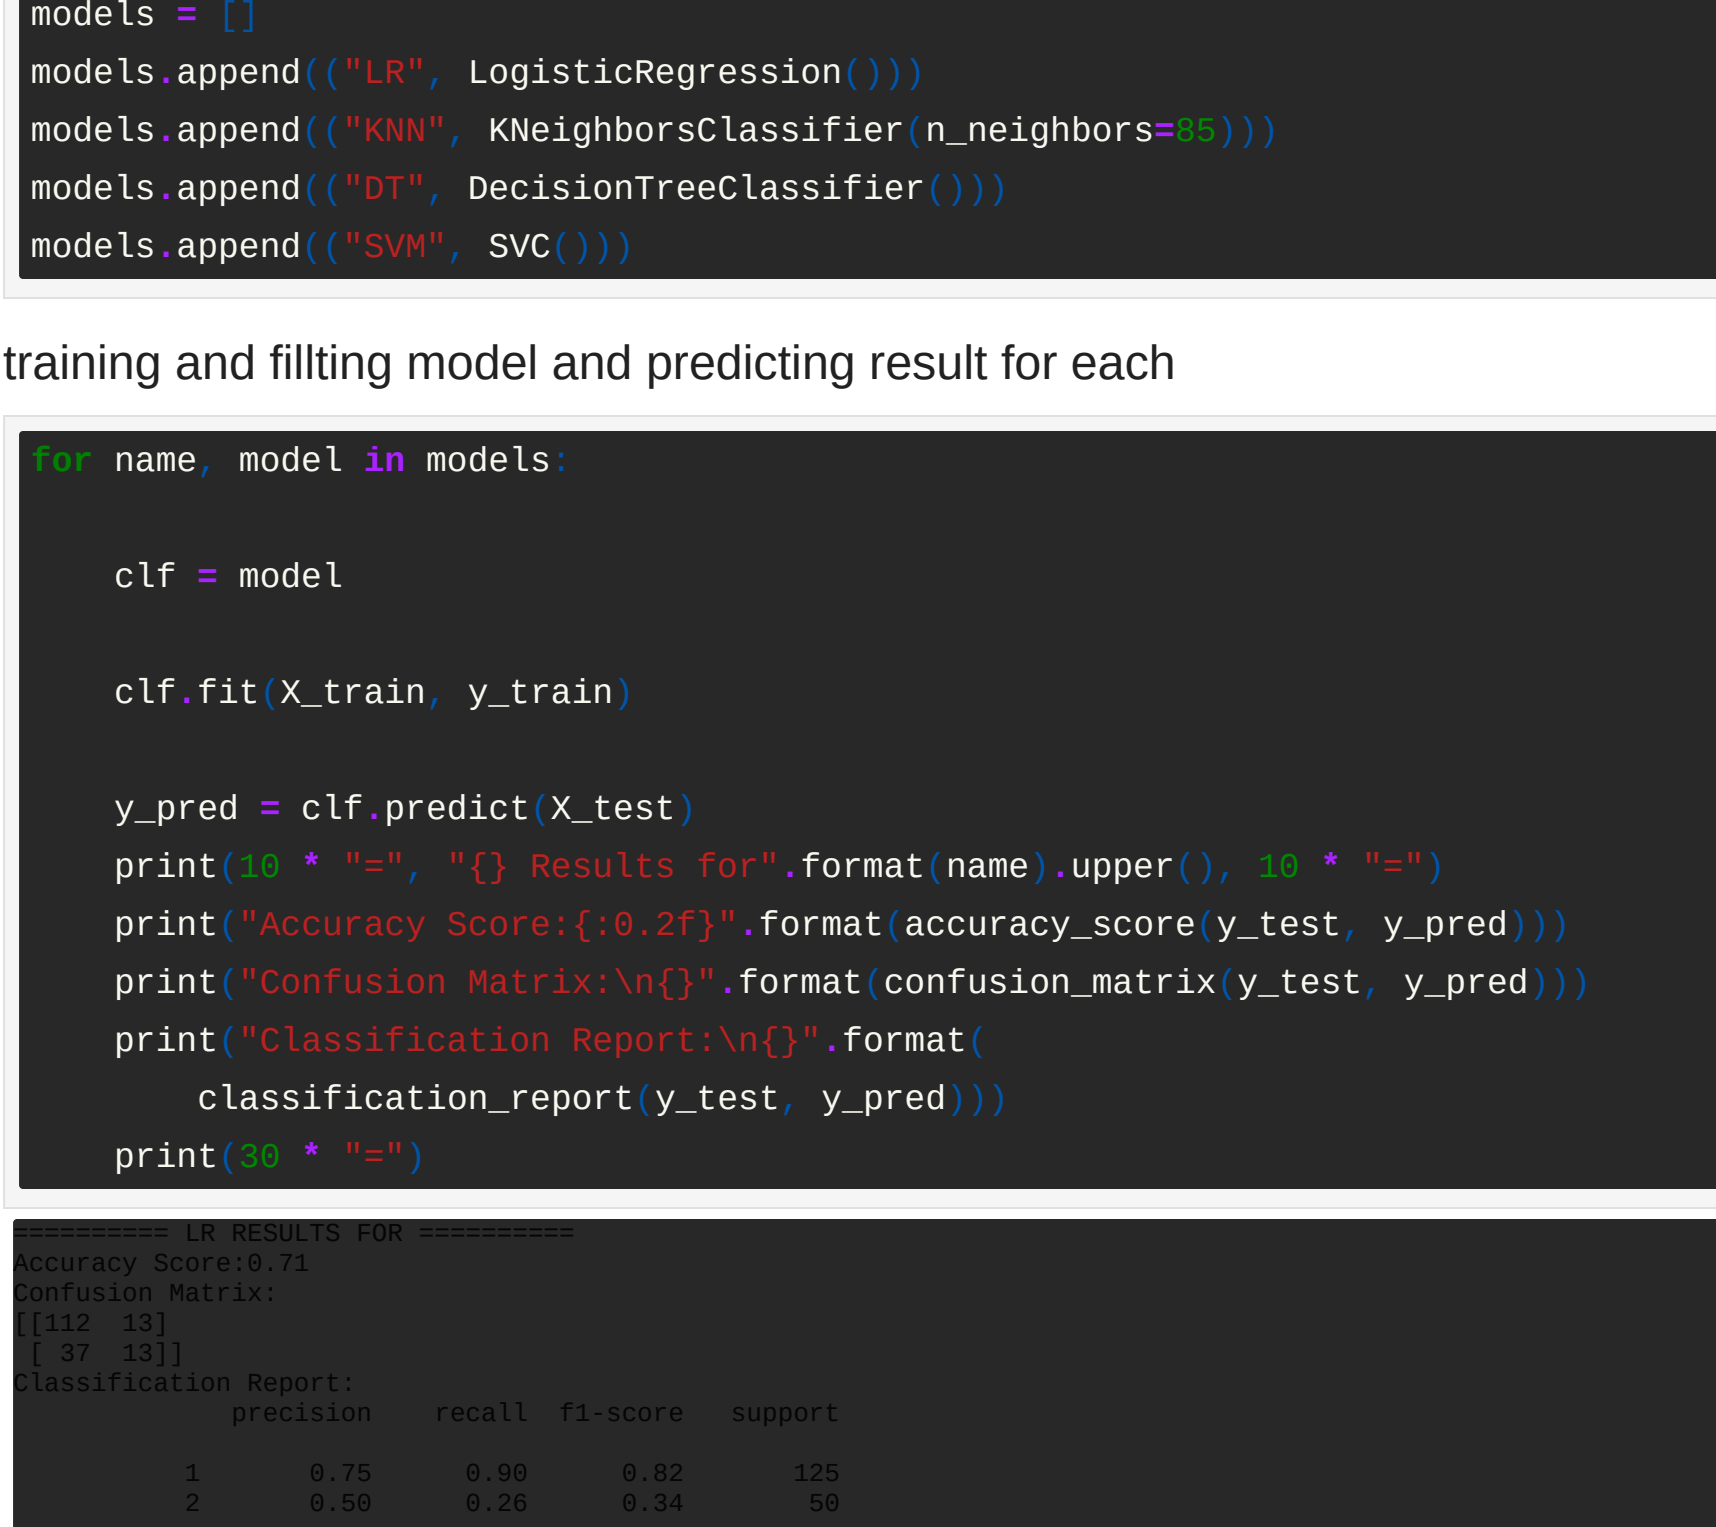
visualizing the correlation between the attributes

In [8]:

```
selected_pair_cols = [
    'Age', 'Total_Bilirubin', 'Direct_Bilirubin', 'Alamine_Aminotransferase',
    'Total_Proteins', 'Dataset'
]

sns.pairplot(data=dataset[selected_pair_cols], hue='Dataset', kind='reg')
```

Out[8]:

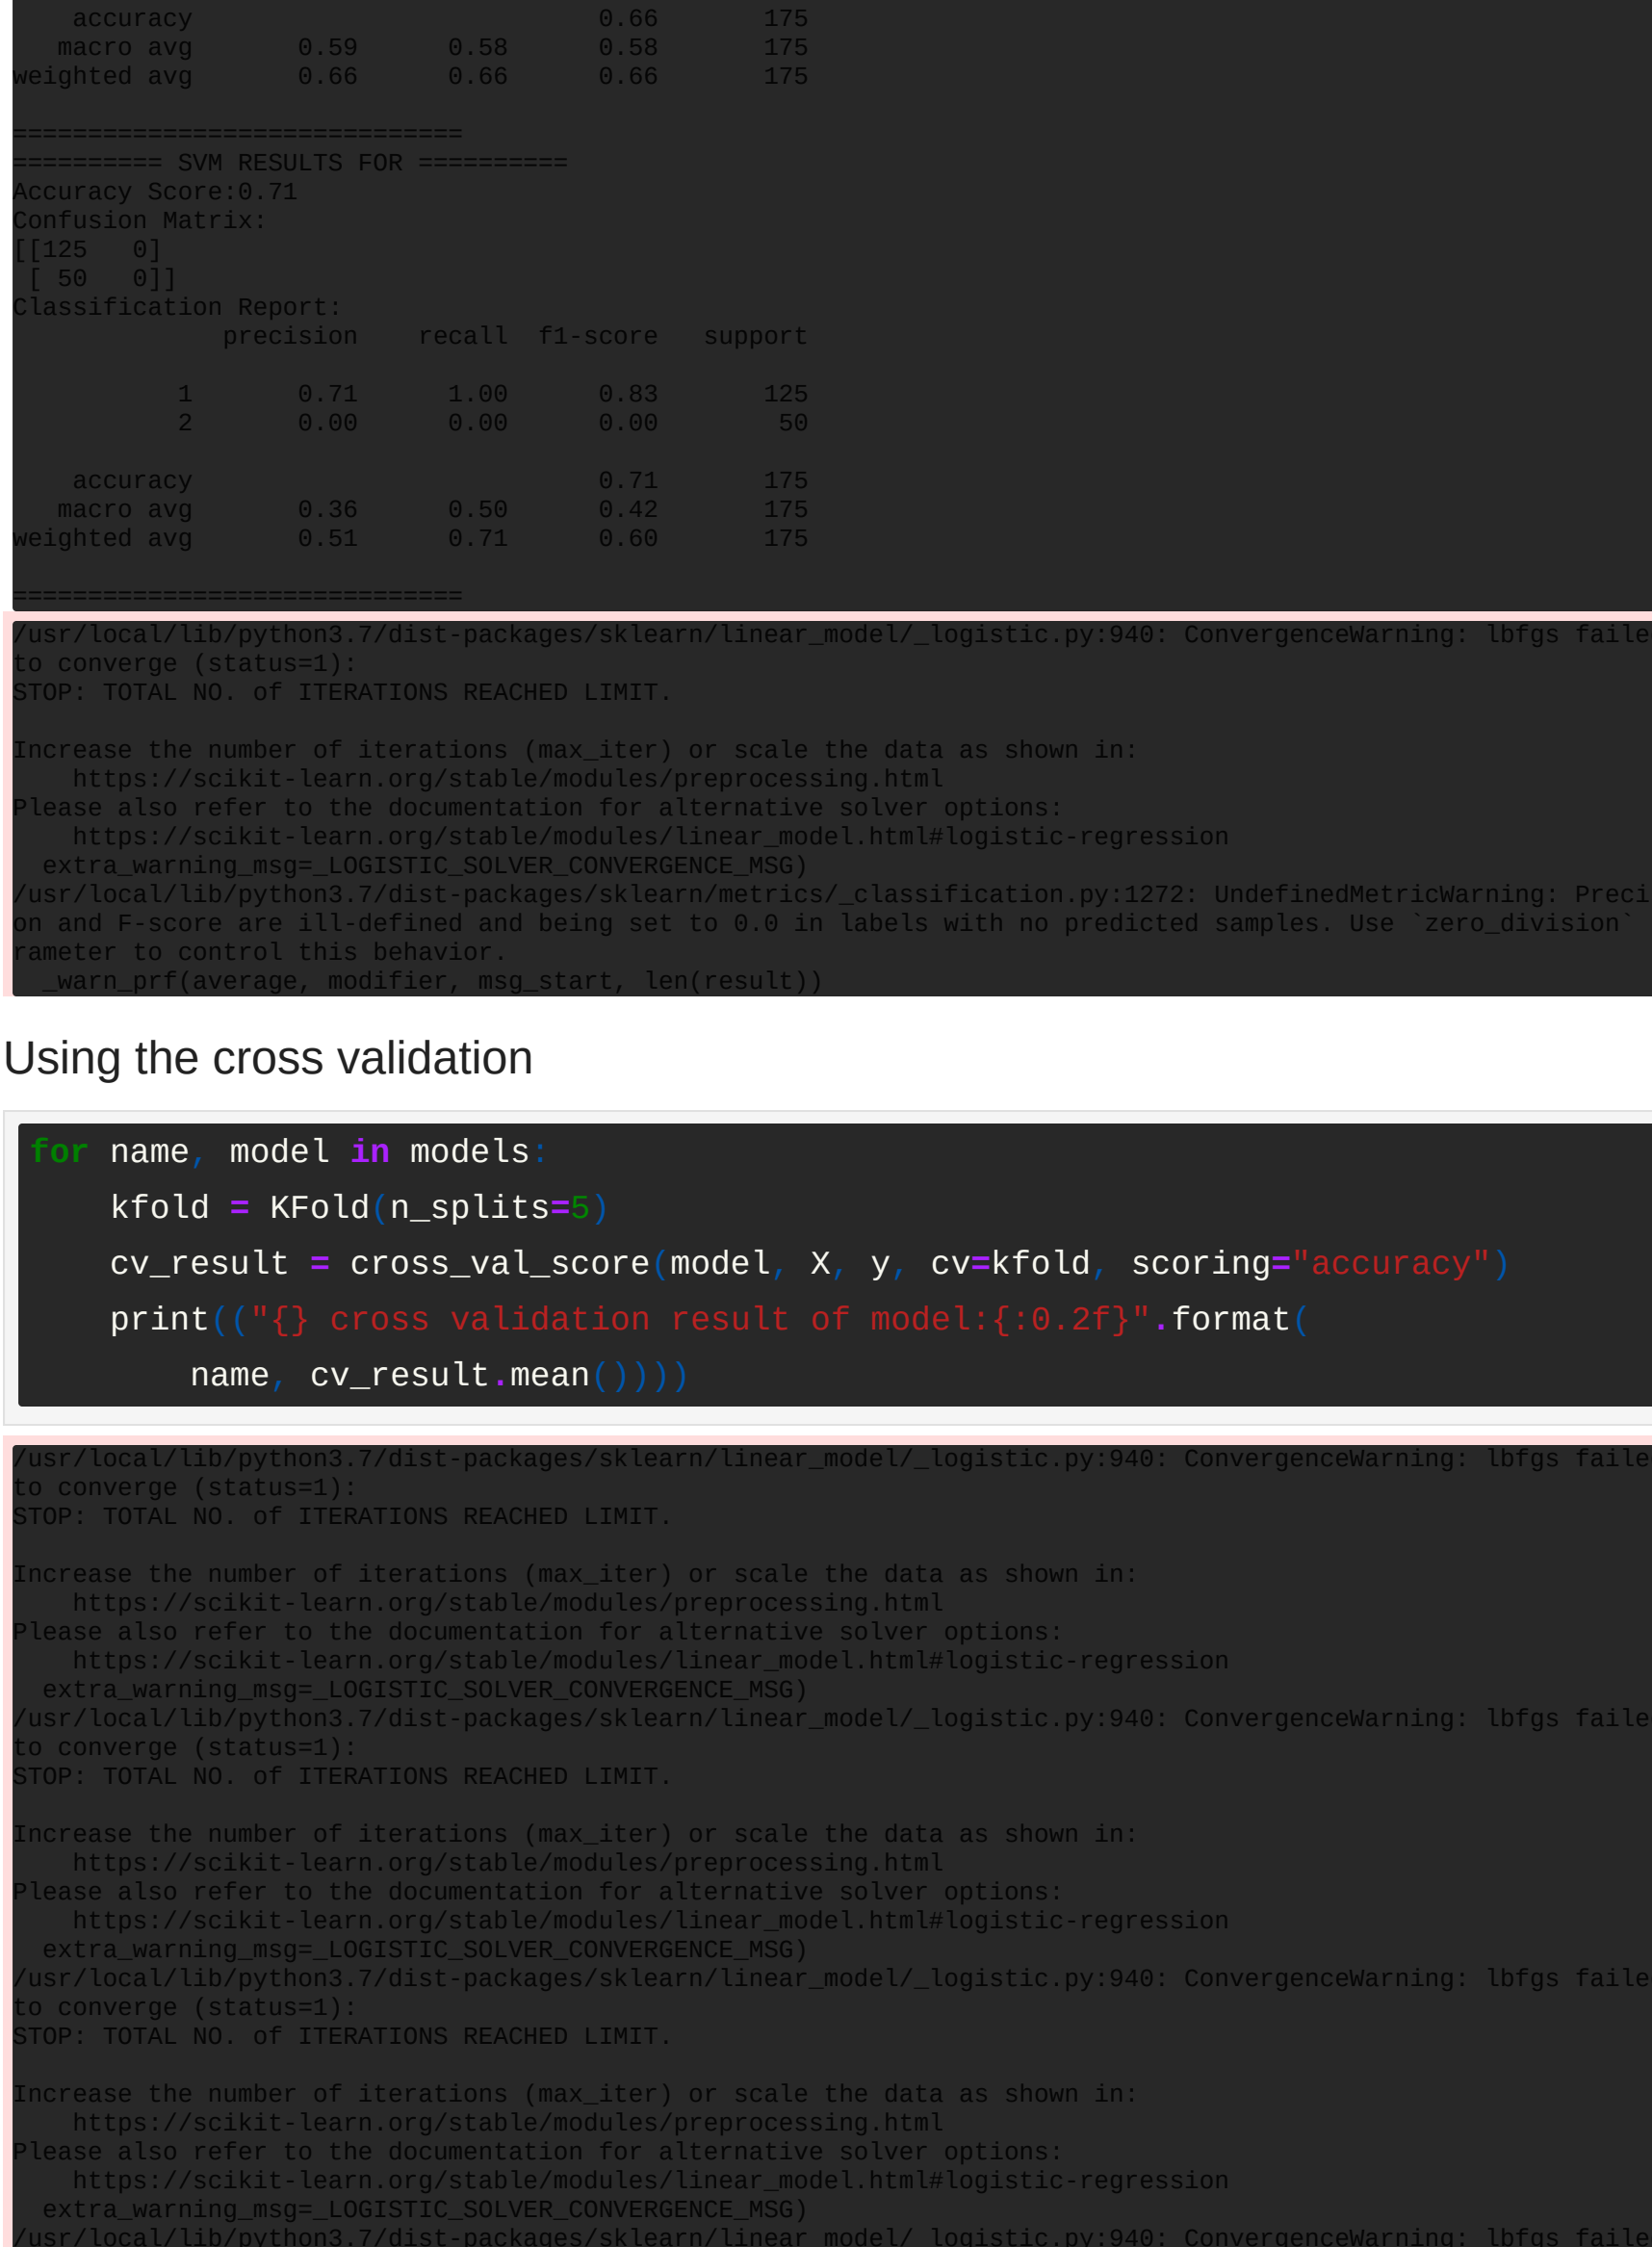


using the correlation matrix

In [9]:

```
corr_matrix = dataset.corr()
fig, ax = plt.subplots(figsize=(15, 15))
sns.heatmap(corr_matrix, annot=True, linewidths=.5, ax=ax)
```

Out[9]:



creating the model

create train and test dataset

In [10]:

```
X = dataset.drop('Dataset', axis=1) # data
y = dataset['Dataset'] # target
```

In [11]:

```
X_train, X_test, y_train, y_test = train_test_split(X,
                                                    y,
                                                    stratify=y,
                                                    test_size=.3,
                                                    random_state=42)
```

list of model to be used

In [12]:

```
models = []

models.append(("LR", LogisticRegression()))
models.append(("KNN", KNeighborsClassifier(n_neighbors= 8)))
models.append(("DT", DecisionTreeClassifier()))
models.append(("SVM", SVC()))
```

training and filling model and predicting result for each

In [13]:

```
for name, model in models:

    clf = model

    clf.fit(X_train, y_train)

    y_pred = clf.predict(X_test)

    print("\n * ==", "\n Results For", format(name, upper), "\n * ==")
    print("Accuracy Score: {:.2f}".format(accuracy_score(y_test, y_pred)))
    print("Confusion Matrix:\n{}".format(confusion_matrix(y_test, y_pred)))
    print("Classification Report:\n{}".format(classification_report(y_test, y_pred)))
    print("\n * ==")

===== LR RESULTS FOR =====
Confusion Matrix:
[[12  13]
 [ 27  15]]
Classification Report:
              precision    recall  f1-score   support

     1         0.76         0.98         0.82         325
     2         0.98         0.25         0.58         58
-
accuracy         0.63         0.58         0.61         375
macro avg         0.66         0.52         0.67         375
weighted avg         0.68         0.71         0.68         375

===== KNN RESULTS FOR =====
Accuracy Score: 0.72
Confusion Matrix:
[[22  21]
 [ 27  15]]
Classification Report:
              precision    recall  f1-score   support

     1         0.72         0.98         0.83         325
     2         0.98         0.25         0.58         58
-
accuracy         0.66         0.52         0.62         375
macro avg         0.69         0.72         0.63         375
weighted avg         0.69         0.72         0.63         375

===== DT RESULTS FOR =====
Accuracy Score: 0.66
Confusion Matrix:
[[26  20]
 [ 27  15]]
Classification Report:
              precision    recall  f1-score   support

     1         0.76         0.77         0.76         325
     2         0.41         0.48         0.44         58
-
accuracy         0.59         0.58         0.66         375
macro avg         0.59         0.59         0.61         375
weighted avg         0.66         0.66         0.66         375

===== SVM RESULTS FOR =====
Accuracy Score: 0.71
Confusion Matrix:
[[25  20]
 [ 27  15]]
Classification Report:
              precision    recall  f1-score   support

     1         0.71         0.98         0.83         325
     2         0.98         0.25         0.58         58
-
accuracy         0.66         0.58         0.62         375
macro avg         0.51         0.71         0.68         375
weighted avg         0.51         0.71         0.68         375

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg: LOGISTIC_SOLVER_CONVERGENCE_MSG)
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg: LOGISTIC_SOLVER_CONVERGENCE_MSG)
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg: LOGISTIC_SOLVER_CONVERGENCE_MSG)
DT cross validation result of model: 0.71
LR cross validation result of model: 0.63
KNN cross validation result of model: 0.72
SVM cross validation result of model: 0.71
```

Using the cross validation

In [14]:

```
for name, model in models:

    kfold = KFold(n_splits=5)

    cv_result = cross_val_score(model, X, y, cv=kfold, scoring='accuracy')

    print("{} cross validation result of model: {:.2f}".format(
        name, cv_result.mean()))

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg: LOGISTIC_SOLVER_CONVERGENCE_MSG)
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg: LOGISTIC_SOLVER_CONVERGENCE_MSG)
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg: LOGISTIC_SOLVER_CONVERGENCE_MSG)
LR cross validation result of model: 0.71
KNN cross validation result of model: 0.72
DT cross validation result of model: 0.63
SVM cross validation result of model: 0.71
```

Finding the Best Classifier with GridSearchCV

In [15]:

```
params_clfs = list()

svm_params = [
    {'kernel': ['rbf'],
     'gamma': [0.5, 1e-3, 1e-4]},
    {'kernel': ['linear'],
     'C': [1, 10, 100, 1000]}
]

params_clfs.append((SVC(), svm_params))

lr_params = {'penalty': ['l1', 'l2'], 'C': np.logspace(-4, 4, 10)}

params_clfs.append((LogisticRegression(), lr_params))

clf = DecisionTreeClassifier()

dt_params = {
    'max_features': ['auto', 'sqrt', 'log2'],
    'min_samples_split': [2, 5, 10, 20, 50, 100, 200, 400, 800, 1600, 3200, 6400],
    'min_samples_leaf': [1],
    'random_state': [100]}

params_clfs.append((DecisionTreeClassifier(), dt_params))

In [16]:
```

```
for clf, param in params_clfs:

    grid_search = GridSearchCV(clf, param, cv=5)

    grid_search.fit(X_train, y_train)

    print("{} * ==", "\n")

    print("{} sklearn result for GridSearch for svm".format(
        clf.__class__.__name__))

    print("{} best params: {}".format(grid_search.best_params_))

    test_means = grid_search.cv_results_['mean_test_score']

    print("{} average test result: {:.2f}".format(np.mean(test_means)))

    y_pred = grid_search.predict(X_test)

    print(
        "{} best parameter result: {:.2f}".format(accuracy_score(y_test,
                                                                y_pred)))

    print("{} Confusion Matrix: \n {}".format(confusion_matrix(y_test, y_pred)))

    plt.figure(figsize=(10, 8))

    sns.heatmap(data=confusion_matrix(y_test, y_pred),
                annot=True,
                cmap="coolwarm",
                xticklabels=['Healthy', 'Unhealthy'],
                yticklabels=['Healthy', 'Unhealthy'])

    plt.xlabel('Predicted values')
    plt.ylabel('Actual values')

    plt.title('Confusion Matrix by Grid Search')

    plt.show()

    print("{} Classification report: \n {}".format(
        classification_report(y_test, y_pred)))

    print("{} * ==", "\n")
```


sklearn result for gridsearch for svm

```
best params: {'C': 30, 'kernel': 'linear'}
```

average test result: 0.72

best parameter result: 0.72

Confusion Matrix:

	Actual \ Predicted	Healthy	Unhealthy
Healthy	124	8	
Unhealthy	41	9	

Confusion Matrix by Grid Search

Classification Report

	precision	recall	f1-score	support
1	0.74	0.84	0.83	125
2	0.53	0.18	0.27	50
accuracy			0.72	175
macro avg	0.63	0.56	0.55	175
weighted avg	0.68	0.72	0.67	175

```

/usr/local/lib/python7.2/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details: ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
/usr/local/lib/python7.2/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
/usr/local/lib/python7.2/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
/usr/local/lib/python7.2/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
/usr/local/lib/python7.2/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
/usr/local/lib/python7.2/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
/usr/local/lib/python7.2/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
/usr/local/lib/python7.2/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

[illegible]

	Healthy	Unhealthy
Healthy	36	1
Unhealthy	14	53

Support values: 0, 10, 20, 30, 40, 50, 60

```

Classification report
precision    recall  f1-score   support

     1       0.72      0.74      0.73      125
     2       0.30      0.28      0.29       50

 accuracy          0.51          0.51          0.51          175
 macro avg          0.51          0.51          0.51          175
 weighted avg          0.48          0.51          0.49          175
  
```

from above we say that classifier gives best result

=====

THE END

=====