

EXPERIMENT NO 3

MULTIPLE LINEAR REGRESSION BY SCRATCH

Importing Modules

```
In [1]: # importing the modules
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import rcParams

# parameters to make the visual look better
rcParams['figure.figsize'] = (14, 7)
rcParams['axes.spines.top'] = False
rcParams['axes.spines.right'] = False
```

Creating The Model Class

```
In [2]: class LinearRegression:
'''
A class which implements linear regression model with gradient descent.
'''
def __init__(self, learning_rate=0.01, n_iterations=10000):
    self.learning_rate = learning_rate
    self.n_iterations = n_iterations
    self.weights, self.bias = None, None
    self.loss = []

'''
function to calculate mean square error
'''

def _mean_squared_error(self, y, y_hat):
'''
Private method, used to evaluate loss at each iteration.

:param y - array, true values
:param y_hat - array, predicted values
:return: float
'''
    error = 0
    for i in range(len(y)):
        error += (y[i] - y_hat[i])**2
    return error / len(y)

'''
function to fit the model
'''

def fit(self, X, y):
'''
Used to calculate the coefficient of the linear regression model.

:param X: array, features
:param y: array, true values
:return: None
'''
    # 1. Initialize weights and bias to zeros
    self.weights = np.zeros(X.shape[1])
    self.bias = 0

    # 2. Perform gradient descent
    for i in range(self.n_iterations):
        # Line equation
        y_hat = np.dot(X, self.weights) + self.bias
        loss = self._mean_squared_error(y, y_hat)
        self.loss.append(loss)

        # Calculate derivatives
        partial_w = (1 / X.shape[0]) * (2 * np.dot(X.T, (y_hat - y)))
        partial_d = (1 / X.shape[0]) * (2 * np.sum(y_hat - y))

        # Update the coefficients
        self.weights -= self.learning_rate * partial_w
        self.bias -= self.learning_rate * partial_d

'''
function to find prediction of the model
'''

def predict(self, X):
'''
Makes predictions using the line equation.

:param X: array, features
:return: array, predictions
'''
    return np.dot(X, self.weights) + self.bias
```

Reading the dataset

```
In [3]: from sklearn.datasets import load_diabetes

data = load_diabetes()
X = data.data
y = data.target
```

Splitting Data Into Train & Test Dataset

```
In [4]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X,
                                                    y,
                                                    test_size=0.2,
                                                    random_state=42)
```

Training the Model

```
In [5]: model = LinearRegression()
model.fit(X_train, y_train)
preds = model.predict(X_test)

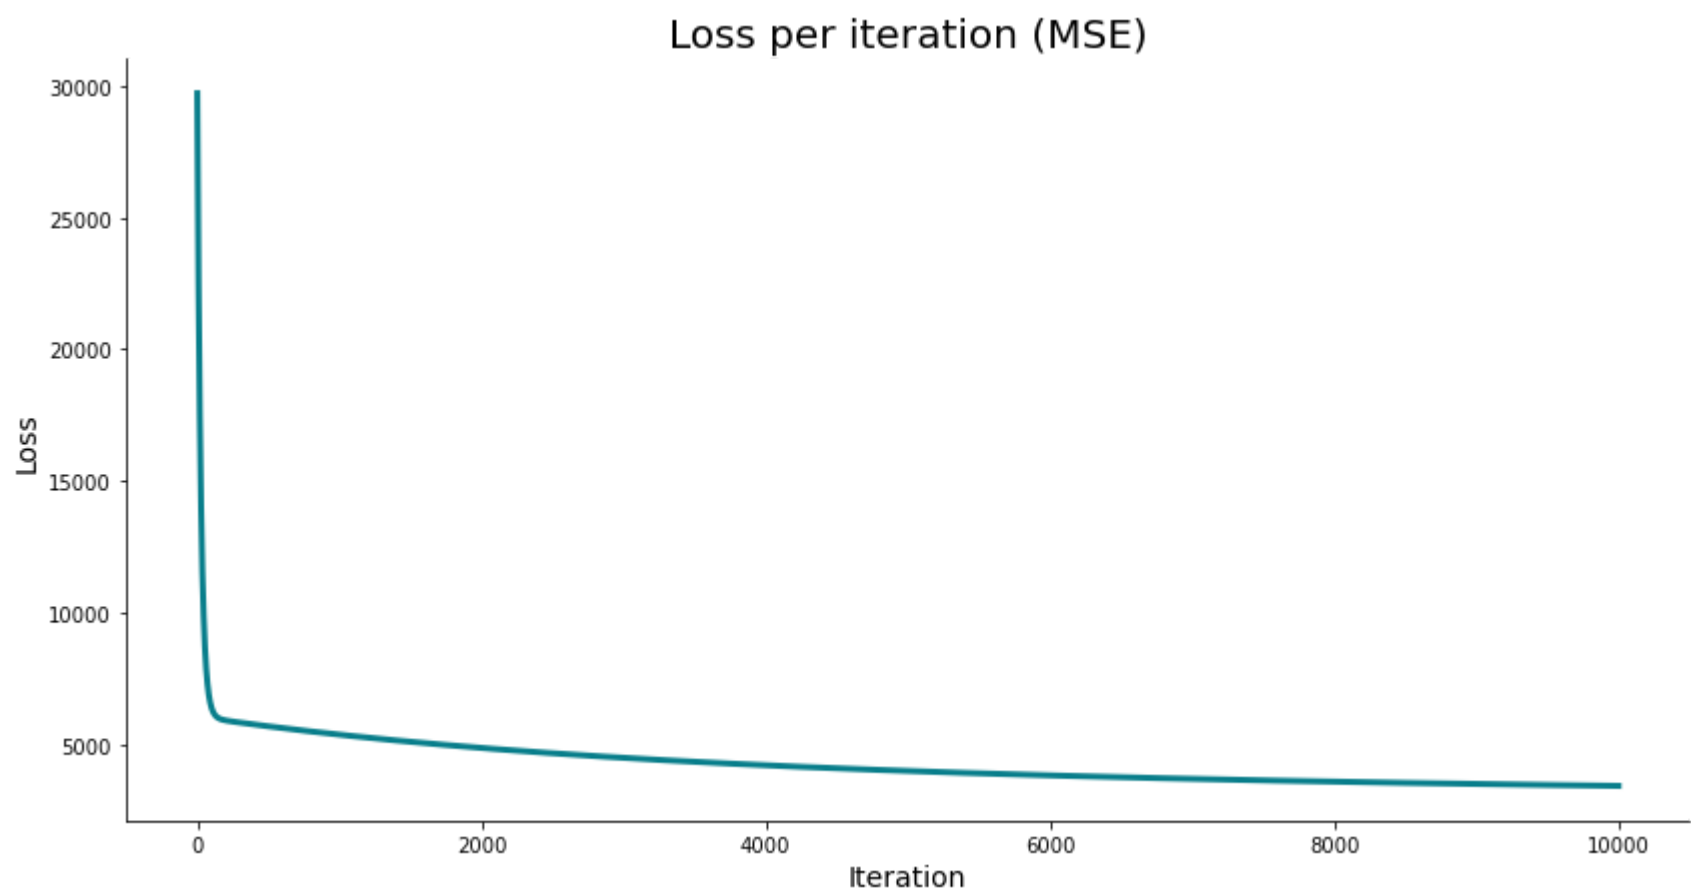
# first few predicted values
preds[:10]
```

```
Out[5]: array([164.21205411, 159.47509229, 161.03082858, 251.53128339,
149.27798242, 131.37880168, 219.01979007, 197.83564132,
114.99909736, 137.90827135])
```

Visualizing of the Loss

```
In [6]: xs = np.arange(len(model.loss))
ys = model.loss

plt.plot(xs, ys, lw=3, c='#087E8B')
plt.title('Loss per iteration (MSE)', size=20)
plt.xlabel('Iteration', size=14)
plt.ylabel('Loss', size=14)
plt.show()
```




Visualizing the Loss at Diffrent Learning Rates

```
In [7]: losses = {}
for lr in [0.5, 0.1, 0.01, 0.001]:
    model = LinearRegression(learning_rate=lr)
    model.fit(X_train, y_train)
    losses[f'LR={str(lr)}'] = model.loss

xs = np.arange(len(model.loss))

plt.plot(xs, losses['LR=0.5'],
         lw=3,
         label=f"LR = 0.5, Final = {losses['LR=0.5'][-1]:.2f}")
plt.plot(xs, losses['LR=0.1'],
         lw=3,
         label=f"LR = 0.1, Final = {losses['LR=0.1'][-1]:.2f}")
plt.plot(xs, losses['LR=0.01'],
         lw=3,
         label=f"LR = 0.01, Final = {losses['LR=0.01'][-1]:.2f}")
plt.plot(xs, losses['LR=0.001'],
         lw=3,
         label=f"LR = 0.001, Final = {losses['LR=0.001'][-1]:.2f}")
plt.title('Loss per iteration (MSE) for different learning rates', size=20)
plt.xlabel('Iteration', size=14)
plt.ylabel('Loss', size=14)
plt.legend()
plt.show()
```



Scoring the Model

```
In [8]: # the learning rate of 0.5 is good so we will select it and train the model again and find the Mean Squared Error
model = LinearRegression(learning_rate=0.5)
model.fit(X_train, y_train)
preds = model.predict(X_test)

model._mean_squared_error(y_test, preds)
```

```
Out[8]: 2885.8431887424276
```

-----X-----X-----X-----X-----X-----X-----X-----

-----X-----X-----