

## Logistic Regression

```
In [ ]: import sklearn
import torch
import torch.nn as nn
import torch.nn.functional as F
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
In [ ]: iris=load_iris()
X=iris['data']
y=iris['target']
```

```
In [ ]: y.shape
```

```
Out[ ]: (150,)
```

```
In [ ]: ### Preprocess the Data
scaler=StandardScaler()
X=scaler.fit_transform(X)
X=torch.tensor(X,dtype=torch.float32)
y=torch.tensor(y,dtype=torch.float32)
```

```
In [ ]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=.2,random_state=32)
```

```
In [ ]: X_train.shape
```

```
Out[ ]: torch.Size([120, 4])
```

```
In [ ]: class LogisticRegression(nn.Module):
    def __init__(self,n_features,h1):
        super(LogisticRegression,self).__init__()
        self.layer1=nn.Linear(n_features,h1)
        self.layer2=nn.Linear(h1,3)
        self.softmax=nn.Softmax(dim=1)

    def forward(self,x):
        out=F.relu(self.layer1(x))
        out=F.relu(self.layer2(out))
        out=self.softmax(out)
        return out
```

```
In [ ]: def train(model,X,y,epocs=50,optimizer=None,loss_fn=None,bsz=20,seed=20):
    torch.manual_seed(seed)
    cost=[]
    for epoch in range(epocs):
        suffle_idx=torch.randperm(y.size(0))
        batches =torch.split(suffle_idx,bsz)
        for batch_idx in batches:
            ###1. Compute the output
            yhat=model(X[batch_idx])

            ###2. Compute the error

            loss=loss_fn(yhat,y[batch_idx].long())

            ###3. Compute the Gradient
            optimizer.zero_grad()
            loss.backward()

            ###4. Update the Parameters
            optimizer.step()


            ### Log the Loss
        with torch.no_grad():
            yhat=model(X)
            curr_loss=loss_fn(yhat,y.long())
            print('Epoc:%3d ' % epoch,end="")
            print('BCE %.2f' % curr_loss)
            cost.append(curr_loss)
    return cost
```

```
In [ ]: model=LogisticRegression(n_features=X_train.size(1),h1=20)
optimizer=torch.optim.Adam(model.parameters(),lr=.001)
loss_fn=nn.CrossEntropyLoss()
```

```
In [ ]: cost =train(model,X_train,y_train,optimizer=optimizer,loss_fn=loss_fn)
```

```
Epoc: 0 BCE 1.13
Epoc: 1 BCE 1.12
Epoc: 2 BCE 1.11
Epoc: 3 BCE 1.10
Epoc: 4 BCE 1.10
Epoc: 5 BCE 1.09
Epoc: 6 BCE 1.09
Epoc: 7 BCE 1.08
Epoc: 8 BCE 1.08
Epoc: 9 BCE 1.07
Epoc: 10 BCE 1.06
Epoc: 11 BCE 1.06
Epoc: 12 BCE 1.05
Epoc: 13 BCE 1.05
Epoc: 14 BCE 1.05
Epoc: 15 BCE 1.04
Epoc: 16 BCE 1.04
Epoc: 17 BCE 1.03
Epoc: 18 BCE 1.03
Epoc: 19 BCE 1.03
Epoc: 20 BCE 1.02
Epoc: 21 BCE 1.02
Epoc: 22 BCE 1.02
Epoc: 23 BCE 1.01
Epoc: 24 BCE 1.01
Epoc: 25 BCE 1.01
Epoc: 26 BCE 1.01
Epoc: 27 BCE 1.01
Epoc: 28 BCE 1.00
Epoc: 29 BCE 1.00
Epoc: 30 BCE 1.00
Epoc: 31 BCE 1.00
Epoc: 32 BCE 1.00
Epoc: 33 BCE 1.00
Epoc: 34 BCE 1.00
Epoc: 35 BCE 1.00
Epoc: 36 BCE 0.99
Epoc: 37 BCE 0.99
Epoc: 38 BCE 0.99
Epoc: 39 BCE 0.99
Epoc: 40 BCE 0.99
Epoc: 41 BCE 0.99
Epoc: 42 BCE 0.99
Epoc: 43 BCE 0.99
Epoc: 44 BCE 0.99
Epoc: 45 BCE 0.99
Epoc: 46 BCE 0.99
Epoc: 47 BCE 0.99
Epoc: 48 BCE 0.99
Epoc: 49 BCE 0.99
```

```
In [ ]: import matplotlib.pyplot as plt
plt.plot(range(len(cost)),cost)
plt.ylabel('Binary Cross Entropy')
plt.xlabel('Epoc')
plt.show()
```



```
In [ ]: from sklearn.metrics import accuracy_score
```

```
In [ ]: with torch.no_grad():
    train_pred=model(X_train)
    #train_pred
    train_acc=(torch.argmax(train_pred,dim=1)==y_train).float().mean()
    print('Training Accuracy: %.2f' %train_acc)
    test_pred=model(X_test)
    test_acc=(torch.argmax(test_pred,dim=1)==y_test).float().mean()
    print('Testing Accuracy: %.2f' %test_acc)
```

Training Accuracy: 0.65

Testing Accuracy: 0.70

## Multinomial Logistic Regression

```
In [ ]: import torch
import torchvision
import torchvision.transforms as transforms
```

```
In [ ]: from matplotlib.transforms import Transform
transform=transforms.ToTensor()

### Data Loading
traindata=torchvision.datasets.MNIST(root='./data',train=True,download=True,transform=transform)
trainloader=torch.utils.data.DataLoader(traindata,batch_size=32,shuffle=False)

testdata=torchvision.datasets.MNIST(root='./data',train=False,download=True,transform=transform)
testloader=torch.utils.data.DataLoader(testdata,batch_size=32,shuffle=False)
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to ./data/MNIST/raw/train-images-idx3-ubyte.gz

Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to ./data/MNIST/raw/train-labels-idx1-ubyte.gz

Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz

Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw


Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz

Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

```
In [ ]: import matplotlib.pyplot as plt
import numpy as np
def imshow(img):
    img=img.numpy()
    img=np.transpose(img,(1,2,0))
    plt.imshow(img)
    plt.show()
```

```
In [ ]: dataiter=iter(trainloader)
images,labels=dataiter.next()
imshow(torchvision.utils.make_grid(images))
print(labels)
```



```
tensor([5, 0, 4, 1, 9, 2, 1, 3, 1, 4, 3, 5, 3, 6, 1, 7, 2, 8, 6, 9, 4, 0, 9, 1,
        1, 2, 4, 3, 2, 7, 3, 8])
```

```
In [ ]: images.view(images.size(0),-1).shape
```

```
Out[ ]: torch.Size([32, 784])
```

```
In [ ]: print(images[0].shape)
```

```
torch.Size([1, 28, 28])
```

```
In [ ]: import torch.nn as nn
import torch.nn.functional as F
class LogisticNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer1=nn.Linear(1*28*28,100)
        self.layer2=nn.Linear(100,30)
        self.layer3=nn.Linear(30,10)

    def forward(self,x):
        x=F.relu(self.layer1(x))
        x=F.relu(self.layer2(x))
        x=self.layer3(x)
        return x

model=LogisticNet()
```

```
In [ ]: model
```

```
Out[ ]: LogisticNet(
  (layer1): Linear(in_features=784, out_features=100, bias=True)
  (layer2): Linear(in_features=100, out_features=30, bias=True)
  (layer3): Linear(in_features=30, out_features=10, bias=True)
)
```

```
In [ ]: import torch.optim as optim
```

```

In [ ]: def train_mnst(model,train_loader,test_loader,EPOCHS=100,lr=.001,seed=32):
    torch.manual_seed(seed)
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    cost=[]
    model.to(device)
    optimizer=optim.SGD(model.parameters(),lr=lr)
    loss_function=nn.CrossEntropyLoss()
    for epoch in range(EPOCHS):

        curr_loss=0.0
        current=0.0
        t_current=0.0
        for i,data in enumerate(train_loader,0):

            ## Get a Batch of Data
            input,target=data[0].to(device),data[1].to(device)
            ###1. Compute the output

            ###Flatten the input
            input=input.view(input.size(0),-1)
            yhat=model(input)

            ###2. Compute the Loss
            loss=loss_function(yhat,target)

            ###3. Compute the gradient

            optimizer.zero_grad()
            loss.backward()

            ###4. Update the parameters

            optimizer.step()
            ### Log the record
            with torch.no_grad():
                yhat=model(input)
                err=loss_function(yhat,target)
                curr_loss+=err
                b_current=((torch.argmax(yhat,dim=1)==target).sum()).float()

                current+=b_current
            ### Calculate the accuracy
            with torch.no_grad():
                for data in test_loader:
                    ims,lbs=data[0].to(device),data[1].to(device)
                    pred=model(ims.view(ims.size(0),-1))
                    t_current+=(torch.argmax(pred,dim=1)==lbs).sum()

        print('Current Epoch: %3d' %(epoch+1),end="")
        print(' | Loss: %5.2f' %curr_loss,end="")
        accuracy=current.sum()/len(traindata)*100
        test_accuracy=t_current.sum()/len(testdata)*100
        #print(current)
        print(' | Accuracy: %.2f' %accuracy,end="")
        print(' | Test Accuracy: %.2f' %test_accuracy )
    return cost

```

```
In [ ]: cost=train_mnst(model,trainloader,testloader,EPOCHS=30)
```

```
-----  
NameError                                Traceback (most recent call last)  
<ipython-input-2-9ab78b465309> in <module>()  
----> 1 cost=train_mnst(model,trainloader,testloader,EPOCHS=30)  
  
NameError: name 'model' is not defined
```

```
In [ ]: len(traindata)
```

```
Out[ ]: 60000
```

```
In [ ]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
prob=model(images.view(32,-1).to(device))  
(torch.argmax(prob,dim=1).to('cpu')==labels).sum().float()/32*100
```

```
Out[ ]: tensor(93.7500)
```